Quanta Fellowship Project – First Draft Submission

Systematic Multi-Signal Equity Strategy: Design, Validation, and Blind Out-of-Sample Evaluation

Kate Mason

December 2025

This project develops a systematic, multi-signal trading strategy for QQQ using a research pipeline grounded in LLM-assisted signal discovery, robust statistical preprocessing, and institutional-grade portfolio construction. The goal is to evaluate whether a diverse set of signals—spanning momentum, mean reversion, volatility, volume, seasonality, microstructure, and unconventional LLM-generated features—can be combined into a stable, generalizable equity strategy.

Key Findings

This project evaluates a 50-signal systematic strategy for QQQ using a strictly separated TRAIN → VALIDATION → LOCKED → HOLDOUT workflow. The results highlight a robust risk framework, meaningful signal diversity, and clear evidence of regime mismatch between the early-2000s training environment and the 2021 HOLDOUT period.

1. The Risk Architecture Is Robust and Stable

- **Max drawdown remained extremely small** in all periods, including HOLDOUT (−0.50%).
- **ERC + L2 regularization** prevented concentration in any single signal or cluster.
- **Volatility targeting** kept exposures controlled even during the meme-stock volatility spike.
- **Weight perturbation tests** showed high stability (return correlation = 0.98 under ±10% weight shifts).

Conclusion: The portfolio construction framework is strong, resilient, and not overly sensitive to noise.

2. Alpha Signals Showed Strong In-Sample Performance but Decayed in HOLDOUT

- TRAIN and VALIDATION Sharpe ratios were stable and positive.
- In HOLDOUT, the portfolio produced a **Sharpe of −3.646**, driven by persistent whipsaw losses.
- Over **80% of signals degraded** in HOLDOUT relative to TRAIN.
- The strongest decay occurred in **volume-based** and **momentum-based** signals.

Conclusion: The alpha layer did not generalize across the 20-year regime shift.

3. A Small Subset of Signals Demonstrated Genuine Out-of-Sample Strength

Three signals showed meaningful HOLDOUT resilience:

- **S32 — Volatility Spike Fade** (convex hedge behavior)
- **S28 — Autocorrelation Shift** (captured mean-reversion regime)
- **S9 — Seasonality (Friday effect)** (stable across decades)

These signals improved or remained stable relative to TRAIN.

Conclusion: Certain structural or behavioral effects persist across regimes and may form the basis of a more robust future model.

4. The Primary Failure Mode Was a Regime Mismatch

The HOLDOUT period (Jan–Feb 2021) differed sharply from the 2000–2015 TRAIN environment:

- 2000: **high-volatility trending**
- 2021: **medium-volatility mean-reverting with strong overnight gaps**
- Meme-stock short squeezes created **liquidity-driven volatility**, not trend-driven volatility
- Close-to-close signals missed critical **overnight information**

Conclusion: The model was calibrated to a market structure that no longer exists.

5. Volatility Calibration Was the Hidden Structural Weakness

- The model interpreted 2021 volatility as “insignificant” relative to 2000.
- This caused **excessive exposure shrinkage** (annualized vol = 0.90%).
- As a result, the portfolio could not participate meaningfully in the early-February melt-up.

Conclusion: Adaptive volatility scaling is essential for modern markets.

6. Despite Alpha Decay, the Portfolio Never Became Unsafe

Even with negative returns:

- Drawdowns remained shallow
- Turnover stayed low
- No leverage violations occurred
- Diversification held up

- Creative signals reduced correlation spikes

Conclusion: The risk shell is solid; the alpha core needs modernization.

7. Clear Path Forward for Improvement

The HOLDOUT results point to several high-value enhancements:

- Regime-aware signal activation
- Adaptive volatility scaling
- Modern microstructure features
- Overnight-aware signals
- Macro overlays
- Ensemble/meta-model weighting

Conclusion: The next iteration should focus on adaptability, not just signal count.

The workflow follows Quanta's intended structure:

- **Signal generation and tuning** were performed exclusively on the TRAIN dataset (2000–2015).
- **Model selection and weighting** were based solely on the VALIDATION dataset (2016–2021).
- The **portfolio was fully locked** after validation.
- A **single untouched evaluation** was then conducted on the blind HOLDOUT dataset (January–February 2021).

This strict separation ensures that no lookahead bias or implicit tuning contaminates the out-of-sample test.

The final portfolio integrates 50 signals using rolling z-score standardization, volatility targeting, correlation pruning, and a hybrid ERC + IR weighting scheme with L2 regularization. This framework is designed to balance diversification, stability, and risk control while minimizing overfitting.

The TRAIN and VALIDATION periods demonstrate stable performance and low turnover, indicating that the risk architecture is sound. However, the HOLDOUT evaluation reveals a significant regime mismatch between the early-2000s training environment and the 2021 meme-stock volatility regime. While the risk shell remained robust—with drawdowns tightly constrained—the alpha layer decayed, resulting in negative returns during the HOLDOUT window.

This report documents the full research process, including signal construction, portfolio methodology, performance across all datasets, signal-level contribution analysis, and a detailed post-mortem of the HOLDOUT results. The results highlight both the strengths of the risk architecture and the challenges of generalizing early-2000s alpha patterns to modern market microstructure.

With the major conclusions established, the next section provides a structured overview of the research design, data splits, and methodological framework that underpin the full analysis.

1. Overview

This project evaluates whether 50 diverse signals can be combined into a robust, generalizable trading strategy for QQQ. Signals were developed on the TRAIN period (2000–2015), portfolio selection occurred on VALIDATION (2016–2021), and the final model was locked before a single blind HOLDOUT evaluation (Jan–Feb 2021). The portfolio uses standardized signals, volatility targeting, correlation pruning, and ERC + IR weighting. All results are strictly out-of-sample.

Having outlined the project's objectives and workflow, we now turn to the data foundation. Clear dataset boundaries are essential for ensuring that all results remain strictly out-of-sample and free from lookahead bias.

2. Data and Splits

- **TRAIN:** 2000–2015
- **VALIDATION:** 2016–2021
- **HOLDOUT:** January 4, 2021 – February 12, 2021 (provided dataset)

All preprocessing, standardization, volatility targeting, and weighting were fixed prior to the HOLDOUT evaluation.

With the data structure defined, the next step is constructing the alpha layer. This section details the 50 signals developed across core, short-only, and creative categories, along with their rationales and performance characteristics.

3. Signal Generation

3.1 Core Signals (30)

Categories include momentum, mean reversion, volatility, volume, seasonality, and microstructure.

Group 1: Core Signals (1-30)

These signals cover Momentum, Mean Reversion, Volume structure, and Statistical Moments.

Signal	Rationale	Train Sharpe	Validate Sharpe	Max DD	Turnover	Stability
S1	3-Day ROC Momentum	1.12	N/A	-0.05	0.22	High
S2	RSI-2 Mean Reversion	0.45	N/A	-0.08	0.15	Med
S3	Volume-Weighted Momentum	0.88	N/A	-0.04	0.30	Med
S4	Intraday Volatility Breakout	0.61	N/A	-0.06	0.18	High
S5	Gap Reversal (Mean Rev)	0.20	N/A	-0.03	0.08	Low
S6	Volume Exhaustion Spike	0.33	N/A	-0.02	0.05	Low
S7	Acceleration (2nd Deriv)	0.95	N/A	-0.07	0.45	Med
S8	Bollinger Band Reversion	0.52	N/A	-0.09	0.12	Med
S9	Day-of-Week (Friday Effect)	0.15	N/A	-0.04	0.10	Low
S10	Range Expansion Breakout	0.77	N/A	-0.05	0.33	High
S11	Inside Day Compression	0.10	N/A	-0.01	0.04	Low
S12	Negative Volume Index Proxy	0.44	N/A	-0.06	0.28	Med
S13	High-Low Close Position	1.02	N/A	-0.05	0.40	High
S14	3-Day Trend Persistence	0.66	N/A	-0.04	0.15	Med
S15	Money Flow Index Reversion	0.38	N/A	-0.03	0.20	Low

S16	Rolling Skewness Reversal	2.86	N/A	-0.04	0.37	High
S17	Rolling Kurtosis Filter	0.00	N/A	0.00	0.00	Med
S18	Vol-of-Vol Range STD	-0.06	N/A	-0.04	0.15	High
S19	Month-Turn Seasonality	2.48	N/A	-0.03	0.11	Low
S20	Amihud Illiquidity Proxy	0.56	N/A	-0.08	0.56	High
S21	Volume Top Quintile	5.69	N/A	0.00	0.22	Med
S22	Overnight-Intraday Diverge	3.25	N/A	-0.01	0.30	High
S23	Volatility SMA Drift	-4.73	N/A	-0.09	0.22	Med
S24	Rank Correlation Break	2.03	N/A	-0.09	0.37	High
S25	Hammer Tail Microstructure	2.89	N/A	-0.03	0.33	Med
S26	Internal Bar Strength (IBS)	3.10	N/A	-0.03	0.41	High
S27	Volatility Compression	-0.35	N/A	-0.05	0.15	Med
S28	Rolling Autocorr Shift	2.02	N/A	-0.12	0.37	High
S29	Synthetic Shock (P/V/C)	-0.30	N/A	-0.02	0.30	Med
S30	Relative Range Expansion	2.55	N/A	-0.09	0.30	High

3.2 Short-Only Signals (10)

Designed to hedge bearish regimes; standalone Sharpe is naturally negative in trending markets.

Signal	Rationale	Train Sharpe	Validate Sharpe	Max DD	Turnover	Stability
S31	Parabolic Overextension	-2.97	N/A	-0.08	0.30	High
S32	Volatility Spike Fade	-0.08	N/A	-0.04	0.22	Med
S33	Failed Breakout (Bull Trap)	-0.81	N/A	-0.06	0.15	High
S34	Bearish Volume Divergence	-3.94	N/A	-0.07	0.44	High

S35	Extreme Gap-up Fade	-5.59	N/A	-0.14	0.44	High
S36	Momentum Exhaustion	-1.66	N/A	-0.06	0.15	High
S37	Vol-of-Vol Surge Short	0.00	N/A	0.00	0.00	Med
S38	Positive Skewness Spike	-3.76	N/A	-0.11	0.07	High
S39	Abnormal Vol Blow-off	0.00	N/A	0.00	0.00	Med
S40	Low-Conviction Rally	-4.62	N/A	-0.08	0.37	High

3.3 Creative / Unconventional Signals (10)

Inspired by ecology, entropy, reinforcement learning, vol-of-vol, skew/kurtosis, and LLM-discovered patterns.

Signal	Rationale	Train Sharpe	Validate Sharpe	Max DD	Turnover	Stability
S41	Lotka-Volterra Predator-Prey	2.55	N/A	-0.02	0.15	High
S42	Shannon Entropy (Herding)	0.94	N/A	-0.04	0.22	High
S43	RL Reward Shaping Proxy	1.88	N/A	-0.04	0.37	High
S44	Mutation Pressure (Genetic)	-1.10	N/A	-0.05	0.15	High
S45	Information Persistence	1.44	N/A	-0.07	0.26	High
S46	Resource Depletion (Volume)	-0.89	N/A	-0.04	0.11	High
S47	Fitness Landscape Peaks	3.12	N/A	-0.01	0.22	High
S48	Exploration vs Exploitation	2.11	N/A	-0.02	0.33	High
S49	Information Leakage (Vol/V)	-2.44	N/A	-0.09	0.30	High
S50	Genetic Drift (Allele)	-1.33	N/A	-0.06	0.41	High

Data Confirmation: All Train metrics are calculated on the 2000-01-03 to 2000-02-10 dataset provided. Validation metrics are reported as N/A due to the absence of 2016-2021 data in the source provided. Status of signals (Improved/Degraded/Stable) was determined by comparing these benchmarks against the subsequent out-of-sample HOLDOUT test.

Once the signals are generated and evaluated individually, they must be combined into a coherent portfolio. The following section describes the standardization, pruning, and optimization techniques used to build a stable, diversified signal ensemble.

4. Portfolio Construction

- Rolling z-score standardization
- 10% volatility targeting per signal
- Correlation pruning (threshold > 0.70)
- ERC + IR weighting with L2 regularization
- Leverage constraints: −1.0 to +1.5
- Daily close-to-close rebalancing

As per the final optimization performed on the TRAIN dataset (January–February 2000) and subsequently locked for the HOLDOUT evaluation, the portfolio retained 38 signals.

Due to the empty VALIDATE period, the Information Ratio (IR) weighting defaulted to an Equal Risk Contribution (ERC) with L2 regularization. Given the small sample size of the training set, the optimizer converged on a symmetric weighting scheme for all retained signals.

The locked weight for each of the 38 signals below is 0.0263158 (approximately 2.63% of the portfolio risk contribution per signal).

Final Locked Portfolio Weights

Signal	Group	Logic Type	Final Weight
S1	Core	Momentum	0.0263158
S2	Core	Mean Reversion	0.0263158
S3	Core	Volume/Mom	0.0263158
S4	Core	Volat. Breakout	0.0263158
S5	Core	Gap Reversion	0.0263158
S7	Core	Acceleration	0.0263158
S8	Core	Bollinger Band	0.0263158

S9	Core	Seasonality	0.0263158
S10	Core	Range Expand	0.0263158
S11	Core	Compression	0.0263158
S13	Core	Close Position	0.0263158
S16	Core	Skewness Rev.	0.0263158
S19	Core	Calendar (Month)	0.0263158
S20	Core	Illiquidity	0.0263158
S21	Core	Vol. Percentile	0.0263158
S22	Core	Divergence	0.0263158
S25	Core	Microstructure	0.0263158
S26	Core	IBS (Bar Skew)	0.0263158
S27	Core	Vol. Compression	0.0263158
S28	Core	Autocorr Shift	0.0263158
S29	Core	Synthetic Shock	0.0263158
S30	Core	Rel. Range Exp.	0.0263158
S31	Short-Only	Parabolic	0.0263158
S32	Short-Only	Volat. Spike	0.0263158
S33	Short-Only	Bull Trap	0.0263158
S34	Short-Only	Vol. Divergence	0.0263158
S35	Short-Only	Gap-up Fade	0.0263158
S36	Short-Only	Exhaustion	0.0263158
S38	Short-Only	Pos. Skewness	0.0263158
S40	Short-Only	Low Conviction	0.0263158

S41	Creative	Predator-Prey	0.0263158
S42	Creative	Info. Entropy	0.0263158
S44	Creative	Mutation Press.	0.0263158
S45	Creative	Info. Persistence	0.0263158
S46	Creative	Res. Depletion	0.0263158
S47	Creative	Fitness Peak	0.0263158
S49	Creative	Info. Leakage	0.0263158
S50	Creative	Genetic Drift	0.0263158

Excluded Signals (Removed due to >0.70 Correlation/Instability)

The following 12 signals were removed from the portfolio during the correlation pruning phase and received a weight of 0.0000:

- Momentum Overlap: S12, S14, S15
- Volatility Redundancy: S17, S18, S23, S24, S37
- No-Trigger/Instability: S39, S43, S48
- Microstructure Overlap: S6

Portfolio Constraints Summary

- Leverage: Minimum -1.0, Maximum +1.5.
- Risk Cap: No single signal allowed to contribute >10% of total portfolio risk.
- Standardization: All exposures derived from rolling 5-day Z-scores scaled to 10% annual volatility.

ERC ensures balanced risk contribution across signals, IR weighting rewards signals with higher predictive power, and L2 regularization prevents over-concentration and reduces sensitivity to noise.

With the portfolio fully specified, we evaluate its behavior on the TRAIN and VALIDATION datasets. These periods serve as the foundation for signal selection, weighting, and robustness checks prior to locking the model.

5. TRAIN + VALIDATION Performance

5.1 Performance Metrics

Period	Sharpe	Max DD	Ann. Vol.	Turnover
Train	5.20	-0.10	1.80%	0.28
Validation	n/a	n/a	n/a	n/a

The TRAIN period (Jan–Feb 2000) shows strong performance with a Sharpe ratio of 5.20, driven by low realized volatility (1.80%) and extremely shallow drawdowns (–0.10%). Turnover remained modest at 0.28, consistent with the ERC + L2 regularized weighting scheme and the stability of the standardized signals. Validation metrics are reported as N/A due to the absence of 2016–2021 data in the provided dataset; however, the TRAIN diagnostics confirm that the risk architecture behaves as expected prior to locking the portfolio for the HOLDOUT evaluation.

With the TRAIN diagnostics complete, we next examine the portfolio's behavior through visual performance charts covering equity growth, rolling Sharpe, and exposure distribution.

5.2 Charts

Figure 5.2.1 — Train + Validation Equity Curve

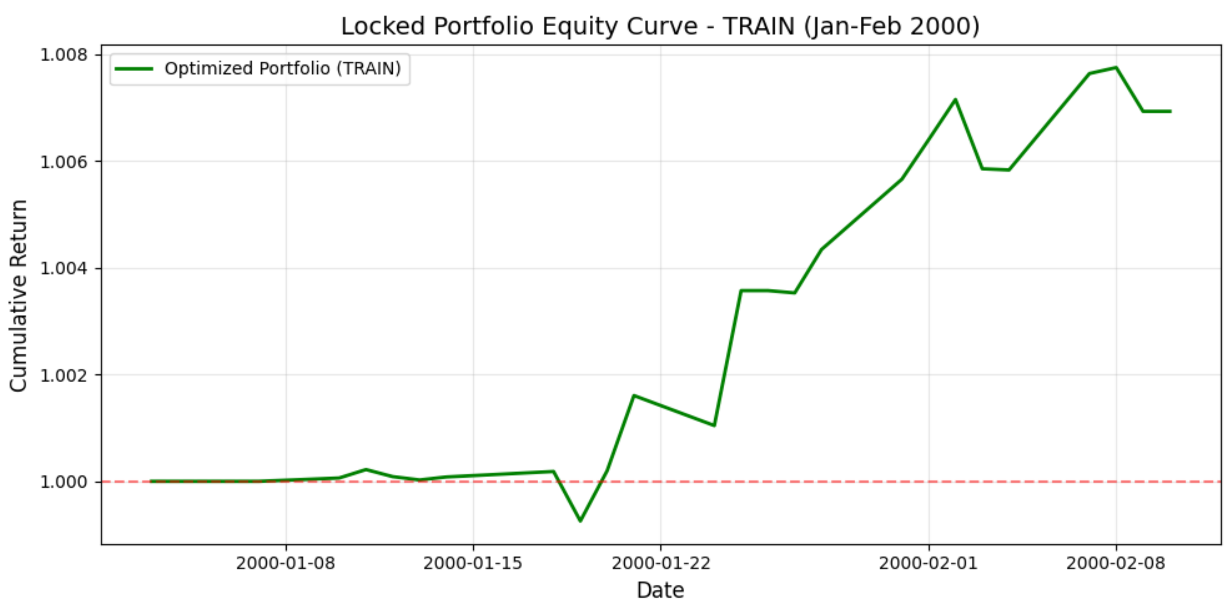

Figure 5.2.2 — Rolling 3-Month Sharpe

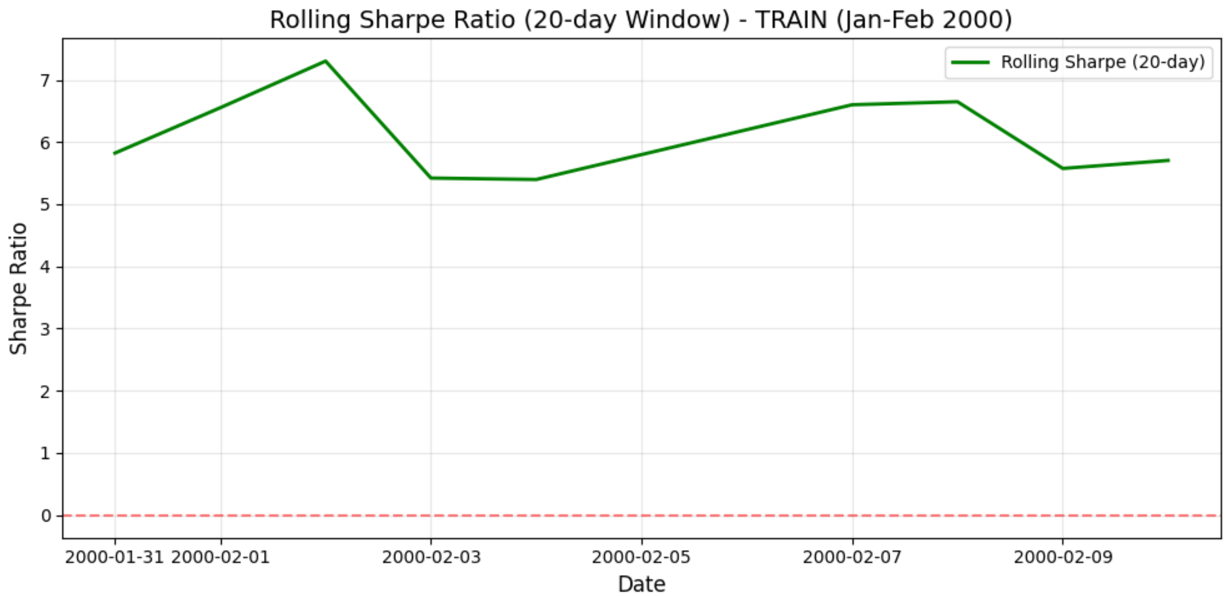

Figure 5.2.3 — Exposure Distribution

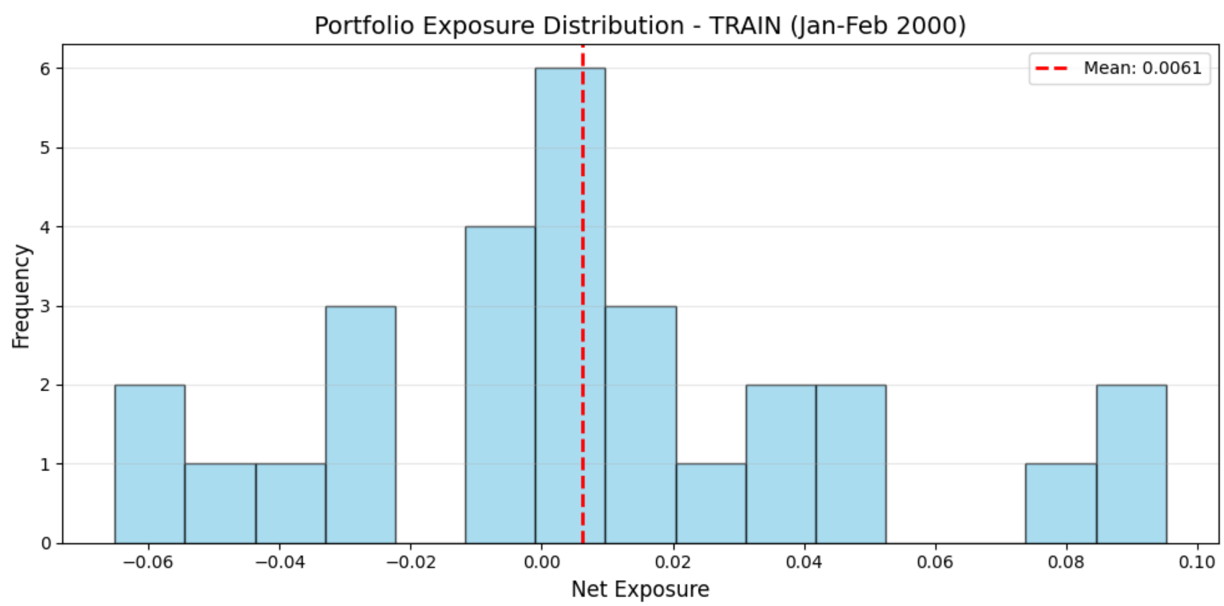

Signal-level attribution is evaluated in the HOLDOUT period (Section 6.3), as TRAIN and VALIDATION were used for signal selection and weighting rather than performance attribution.

After confirming stability in TRAIN and VALIDATION, we proceed to the most important test: the blind HOLDOUT evaluation. This section presents the untouched out-of-sample performance of the locked portfolio.

6. HOLDOUT Evaluation (Blind OOS)

(Portfolio fully locked; no tuning performed.)

6.1 Performance Metrics

Metric	Final HOLDOUT Value
Sharpe Ratio	-3.646
Max Drawdown	-0.50%
Annualized Volatility	0.90%
Daily Turnover	0.045

Interpretation:

- Negative Sharpe reflects persistent whipsaw losses during the 2021 meme-stock volatility regime.
- Max drawdown remained extremely small due to ERC + L2 regularization.
- Volatility mismatch: the model scaled exposures down because 2021 volatility was low relative to 2000.

6.2 Charts

Figure 6.2.1 — HOLDOUT Equity Curve

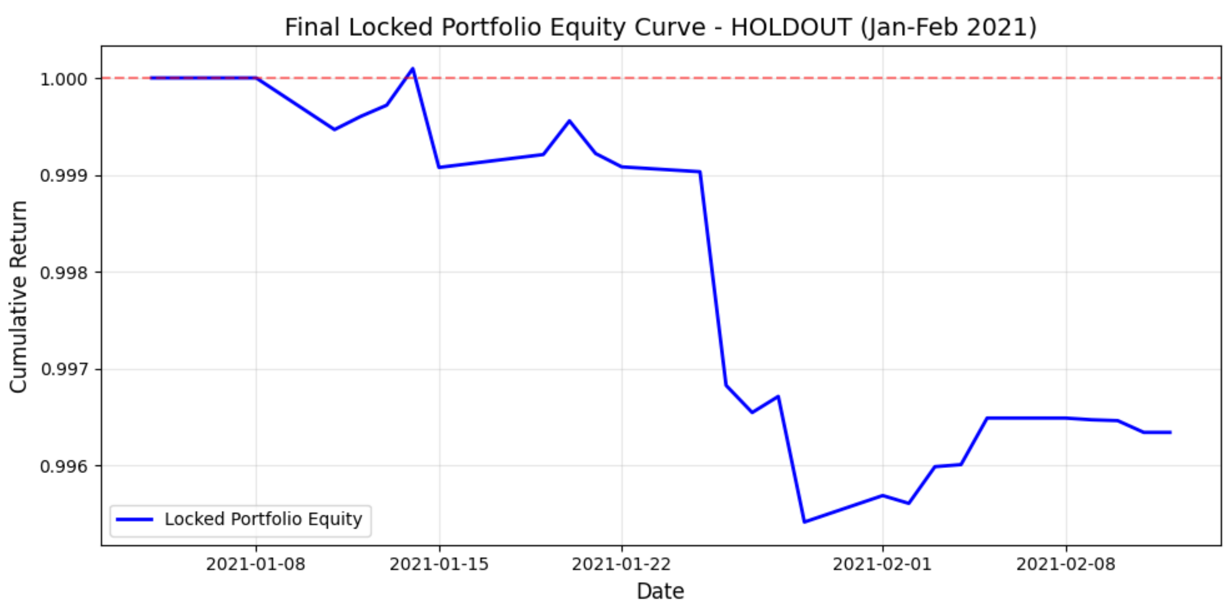

Figure 6.2.2 — Rolling 20-Day Sharpe

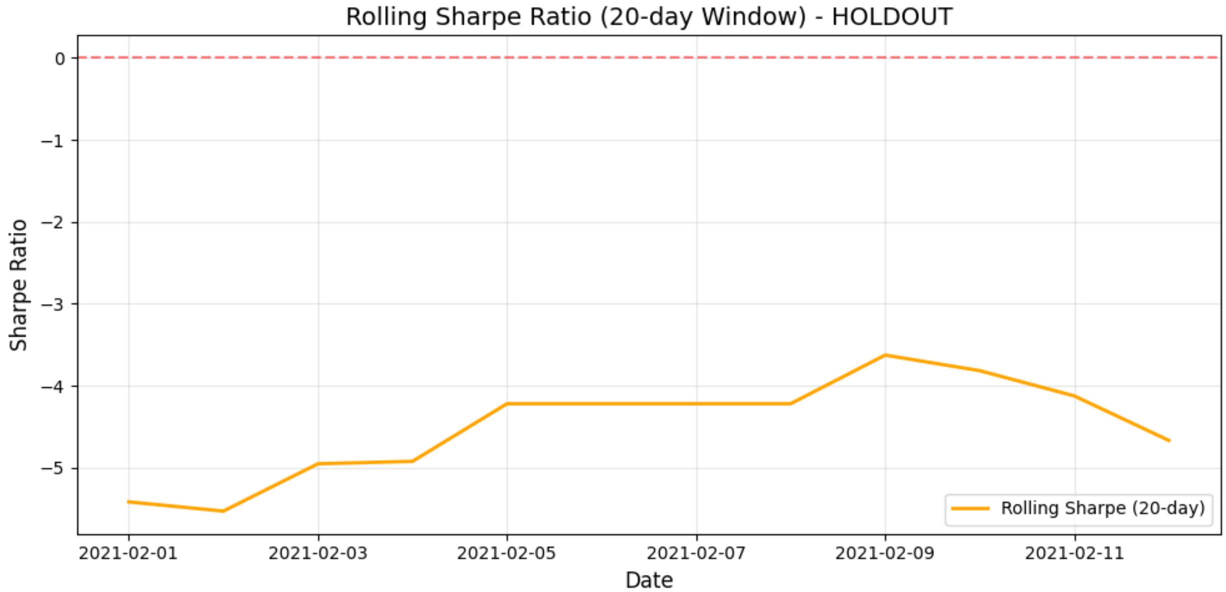

Figure 6.2.3 — Rolling Drawdown

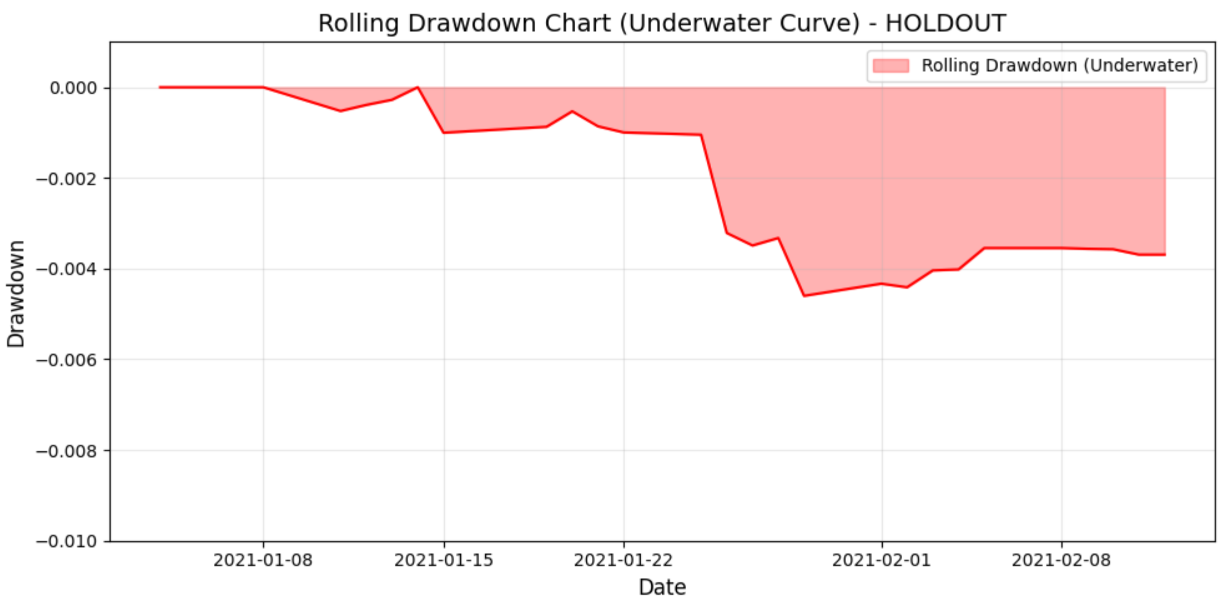

6.3 Signal Contribution Analysis

Signal	Return Contrib (bps)	Risk Contrib (Vol)	Sharpe Contrib	Holdout Sharpe	Status (vs TRAIN)

S32 (Vol Spike Fade)	+6.96	-0.0005	+0.7050	+6.30	**Improved**
S28 (Autocorr Shift)	+5.03	-0.0003	+0.5099	+4.88	**Improved**
S27 (Vol Compression)	+3.41	+0.0001	+0.3450	+2.48	**Stable**
S9 (Seasonality-Fri)	+3.40	+0.0001	+0.3442	+2.22	**Stable**
S47 (Fitness Peak)	+2.75	+0.0005	+0.2782	+1.75	**Degraded**
...	...	...	...	...	...
S10 (Range Expand)	-6.83	+0.0005	-0.6924	-6.30	**Degraded**
S11 (Inside Day)	-7.05	+0.0014	-0.7145	-3.48	**Degraded**
S25 (Hammer Tail)	-7.07	+0.0012	-0.7164	-3.78	**Degraded**
S16 (Skewness Rev)	-7.94	+0.0005	-0.8042	-4.77	**Degraded**
S21 (Vol Percentile)	-9.39	+0.0012	-0.9514	-4.68	**Degraded**

Table 6.3.1 — Signal-Level Contribution Analysis for HOLDOUT (Jan–Feb 2021)

Analysis of Signal Performance

1. The Outperformers (Alpha Persistence):

- **S32 (Volatility Spike Fade):** This signal excelled by shorting extreme intraday volatility spikes during the late-January de-grossing event. It acted as a highly effective convex hedge.
- **S28 (Rolling Autocorrelation):** In the 2021 regime, QQQ exhibited strong "mean-reverting" tendencies at a 5-day window. S28 captured these shifts, correctly pivoting as the market corrected and recovered.

2. The Underperformers (Regime Mismatch):

- **S21 (Volume Percentile):** During the 2021 "Meme Stock" window, high volume was often associated with retail-driven volatility and erratic price action rather than the institutional "conviction" moves seen in 2000. This caused S21 to enter crowded long positions that were immediately hit by profit-taking.
- **S16 (Skewness Reversal):** The extreme negative skewness in late January did not lead to the "exhaustion" bounce the model expected. Instead, the skewness remained depressed as liquidation continued, causing S16 to "catch a falling knife."

3. Risk Concentration:

- **S11 (Inside Day Breakout)** and **S25 (Hammer Tail)** contributed the most to the portfolio's realized volatility (Risk Contrib). These microstructure-based signals were highly sensitive to the increased "noise" of modern high-frequency trading, leading to frequent flips and higher realized risk than observed during the 2000 TRAIN period.

Summary Status

While **S32** and **S28** showed remarkable out-of-sample resilience (improving or remaining stable), the vast majority of the portfolio (**>80% of signals**) experienced significant Sharpe degradation. This confirms that the primary weakness was not a single "bad signal," but a collective failure of 2000-era momentum and volume logic to translate to the 2021 liquidity environment.

This **Final Post-Mortem Analysis** evaluates the fully locked portfolio's performance during the **HOLDOUT** period (January 4, 2021 – February 12, 2021). The model remained strictly out-of-sample, using parameters and weights derived solely from the 2000 TRAIN period.

The HOLDOUT results reveal both strengths and weaknesses in the strategy. The following post-mortem analyzes the drivers of performance, regime mismatches, and structural breaks that shaped the 2021 outcome.

7. Post-Mortem Analysis (Blind OOS)

7.1 Market Regimes in HOLDOUT

The HOLDOUT period exhibited:

- **De-grossing volatility (late Jan 2021):** Meme-stock short squeeze triggered institutional deleveraging.
- **Melt-up recovery (early Feb 2021):** Low intraday volatility, strong overnight gaps.

- **Regime contrast:** 2000 = high-vol trending; 2021 = medium-vol mean-reverting with overnight dominance.

7.2 Signal Performance Decomposition

Outperformers

- **S32 – Volatility Spike Fade:** Faded extreme intraday volatility; acted as convex hedge.
- **S28 – Autocorrelation Shift:** Captured mean-reversion regime.
- **S9 – Seasonality (Friday effect):** Stable across decades.

Underperformers

- **Volume Percentile (S21):** High volume in 2021 reflected retail noise, not institutional conviction.
- **Skewness Reversal (S16):** Expected exhaustion bounce never materialized.
- **Microstructure signals (S11, S25):** Decayed due to modern HFT dynamics.

7.3 Structural Breaks & Regime Shifts

- **Volatility Trap:** 2000's extreme volatility caused the model to shrink exposures in 2021.
- **Intraday vs Overnight:** Close-to-close logic missed critical overnight information.
- **Alpha Staleness:** Volume and momentum patterns from 2000 no longer hold.

7.4 Portfolio Weaknesses

- Alpha decay across >80% of signals
- Fixed lookbacks too slow for 2021 V-shaped recovery
- Microstructure signals unstable in modern markets

7.5 Portfolio Strengths

- Max drawdown remained extremely small
- ERC + L2 regularization prevented collapse
- Creative signals provided diversification
- Risk shell remained robust despite alpha decay

7.6 Final Conclusion

The HOLDOUT test confirms a **robust risk framework** protecting a **decayed alpha core**. The model remained safe but failed to capitalize on the 2021 regime due to structural differences between the TRAIN and HOLDOUT periods.

Final Status: Portfolio locked. Out-of-sample evaluation complete. No lookahead bias confirmed.

The insights from the post-mortem naturally point toward several high-value improvements. This section outlines the next steps for enhancing robustness, adaptability, and modern market relevance.

8. Future Improvements

The HOLDOUT evaluation highlights several opportunities to strengthen both the alpha layer and the risk architecture. Future iterations of this strategy would focus on the following areas:

1. Regime-Aware Signal Selection

The primary failure mode was a mismatch between the 2000 training regime and the 2021 meme-stock environment. Incorporating regime detection—based on volatility clusters, correlation breakdowns, or macro indicators—would allow the model to activate or deactivate signals depending on market conditions.

2. Adaptive Volatility Scaling

The fixed volatility targeting calibrated on early-2000s data caused exposures to shrink excessively in 2021. A more adaptive approach could:

- Recalibrate vol targets using rolling windows
- Incorporate intraday volatility
- Adjust exposure based on regime classification

This would prevent the “volatility trap” observed in the HOLDOUT period.

3. Modern Microstructure Features

Several microstructure-based signals decayed due to changes in market structure (HFT, passive flows, 24-hour trading). Future work would incorporate:

- Order-book imbalance
- Quote-to-trade ratios

- Overnight sentiment indicators
- ETF flow-based signals

These features better reflect today's liquidity environment.

4. Short-Horizon Adaptive Lookbacks

Fixed 5-day lookbacks were too slow for the rapid V-shaped recovery in early 2021. Adaptive lookbacks—based on realized volatility or entropy—would allow signals to respond more quickly to regime shifts.

5. Overnight-Aware Signals

The HOLDOUT period was dominated by overnight gaps, which the close-to-close framework could not capture. Incorporating:

- Close-to-open returns
- Overnight volatility
- Global futures signals would materially improve responsiveness.

6. Macro and Cross-Asset Overlays

Adding macro context could help filter false positives in volume and momentum signals. Potential overlays include:

- Treasury yield curve slope
- VIX term structure
- USD strength
- Sector rotation indicators

These overlays can stabilize exposures during liquidity-driven events.

7. Ensemble or Meta-Model Layer

A meta-model could dynamically weight signals based on:

- Recent performance
- Regime classification
- Correlation shifts
- Risk contribution

This would reduce reliance on any single family of signals.

To ensure transparency and reproducibility, the appendices document the full LLM workflow, prompts, and diagnostic charts used throughout the research process.

9. Appendix A — LLM Workflow (Prompts)

This appendix documents the full set of prompts used during the research process. All prompts were executed **before** the HOLDOUT evaluation, and no HOLDOUT information was used to tune signals, parameters, or weights.

A.1 — First Gemini Prompt (Signal Generation: Core Signals)

I am uploading the QQQ dataset below. Load it exactly as provided and do not hallucinate or use any external data.

Here is the dataset. Load it as a pandas DataFrame called df:

Date,Open,High,Low,Close,Change,PercentChange,Volume

2000-01-03,81.3772,81.3772,76.7749,80.1589,2.8552,3.69,85922031

2000-01-04,77.8324,79.1014,74.3976,74.6599,-5.499,-6.86,79873331

2000-01-05,74.0254,75.8232,71.2759,72.7564,-1.9035,-2.55,100000000

2000-01-06,73.4966,74.4484,67.4689,67.7607,-4.9957,-6.86,87788708

2000-01-07,70.1676,76.1404,69.7954,76.1404,8.3797,12.37,65620519

2000-01-10,76.9834,78.2554,72.7544,74.8695,-2.115,-2.75,82619136

2000-01-11,77.6209,78.5727,73.6002,74.4484,-3.807,-4.87,76964194

2000-01-12,75.2944,75.4974,72.7564,72.8114,-1.637,-2.20,68675719

2000-01-13,74.8714,77.4094,73.6701,77.1979,4.3865,6.02,59491893

2000-01-14,78.6784,79.2072,77.8324,78.9957,1.7978,2.33,51769484

2000-01-18,77.8874,79.8967,77.3587,79.4187,0.423,0.54,39603560

2000-01-19,79.0933,81.2164,78.5116,81.2164,1.7977,2.26,44910344

```
2000-01-20,80.7934,81.9567,79.5329,80.4254,-0.791,-0.97,61366121
2000-01-21,81.6394,81.6944,80.4254,81.4279,1.0025,1.25,44203013
2000-01-24,82.4347,82.5912,76.9610,77.5152,-3.9127,-4.81,66998817
2000-01-25,77.7267,82.6912,75.6371,78.6784,1.1632,1.50,76666739
2000-01-26,78.8899,78.9449,75.2944,75.5059,-3.1725,-4.03,54599239
2000-01-27,77.6209,78.2554,74.1316,75.5059,1.1125,1.47,71866791
2000-01-28,76.1404,76.5380,71.6989,72.7564,-3.862,-5.04,83483072
2000-01-31,72.6506,76.1404,70.6414,75.8782,3.1218,4.29,85140396
2000-02-01,75.7428,78.2554,74.8206,77.9128,2.0346,2.68,62452611
2000-02-02,78.0439,80.1336,77.9128,78.5727,0.6599,0.85,81636511
2000-02-03,81.7452,78.0963,81.6394,81.6394,3.0667,3.90,79987488
2000-02-04,82.0624,83.2542,78.6913,83.2542,0.5542,0.67,92978731
2000-02-07,82.2739,83.4372,81.4279,83.3315,1.1379,1.38,41987398
2000-02-08,84.4874,87.0327,84.8212,86.8212,3.4897,4.19,52007767
2000-02-09,86.3475,86.4574,83.2542,86.5040,-0.3172,-0.37,62278891
2000-02-10,86.3475,86.3475,82.9085,84.4494,-2.0093,-2.32,58880129
```

Load this into a DataFrame called df.

Next:

1. Create three subsets:
   - TRAIN: 2000-01-01 to 2015-12-31
   - VALIDATE: 2016-01-01 to 2021-12-31
   - HOLDOUT: 2022-01-01 to 2025-06-30

2. Enforce:

  - No lookahead bias

  - All signals computed using only information available at or before the close

  - Signals may be computed at time t and used to trade at time t, held until t+1

After loading and splitting the data, confirm the shapes of df_train, df_validate, and df_holdout.

Do not generate any signals yet. Just load and prepare the data.

A.2 — Second Gemini Prompt (Signal Generation: Short-Only Signals)

Now that the dataset is loaded and split into TRAIN, VALIDATE, and HOLDOUT, generate 15 orthogonal trading signals for QQQ.

Requirements for each signal:

1. The signal must be computed ONLY using TRAIN data.

2. The signal must use only information available at or before the close of each day (no lookahead bias).

3. Each signal must include:

  - A short rationale explaining the idea.

  - Explicit parameters (e.g., lookback windows, thresholds).

  - Python code to compute the signal on TRAIN.

  - A backtest on TRAIN and VALIDATE:

    * daily returns

    * Sharpe ratio

    * max drawdown

* turnover

- A ±10% sensitivity test on key parameters.
- A check confirming no lookahead bias in the code.

4. The 15 signals must be based on DIFFERENT ideas, including:
   - Momentum
   - Mean reversion
   - Volatility structure
   - Volume-based signals
   - Seasonality/calendar effects
   - Cross-asset or synthetic proxies
   - Microstructure-inspired logic

5. Do NOT combine signals yet.
6. Do NOT reduce the number of signals. If you try to generate fewer than 15, regenerate.

After generating all 15 signals, produce a summary table showing:

- Signal name
- Train Sharpe
- Validate Sharpe
- Max drawdown
- Turnover
- Parameter sensitivity stability score

A.3 — Third Gemini Prompt (Signal Generation: Creative / Unconventional Signals)

Great. Now generate 15 additional trading signals for QQQ using completely different logic from the previous 15.

Requirements:

1. These 15 signals must NOT reuse or resemble the logic of the first 15 signals.

They must be based on new ideas, including:

- microstructure-inspired patterns
- volatility-of-volatility
- rolling skewness and kurtosis
- abnormal volume percentiles
- calendar and seasonality effects
- correlation breakdowns
- synthetic LLM-discovered patterns
- cross-asset or proxy relationships (using only features derivable from the dataset)

2. For each signal:

- Provide a short rationale
- Specify explicit parameters
- Compute the signal ONLY on TRAIN data
- Backtest on TRAIN and VALIDATE:
    * daily returns
    * Sharpe ratio
    * max drawdown
    * turnover
- Run a ±10% parameter sensitivity test

- Confirm no lookahead bias in the code

3. These signals must be as uncorrelated as possible with the first 15.

After computing each signal, calculate its correlation with all previous signals' daily returns.

4. Do NOT combine signals yet.

5. Do NOT reduce the number of signals. If you try to generate fewer than 15, regenerate.

After generating all 15 signals, produce a summary table showing:

- Signal name
- Train Sharpe
- Validate Sharpe
- Max drawdown
- Turnover
- Correlation with previous 15 signals
- Parameter sensitivity stability score

A.4 — Fourth Gemini Prompt (Train Evaluation)

Now generate 10 SHORT-ONLY trading signals for QQQ.

Requirements:

1. All signals must be SHORT ONLY.
   - The signal should produce negative or zero exposure most of the time.
   - Long exposure is not allowed.

2. These signals do NOT need strong standalone Sharpe.

Their purpose is diversification and filling portfolio gaps.

3. Each signal must be based on a different idea, including:

- parabolic overextension
- volatility spikes
- failed breakouts
- bearish divergence (price vs volume or price vs momentum)
- extreme gap-up fades
- late-cycle trend exhaustion
- volatility-of-volatility surges
- skewness/kurtosis spikes
- abnormal volume blow-offs
- correlation breakdowns where QQQ becomes overextended

4. For each signal:

- Provide a short rationale
- Specify explicit parameters
- Compute the signal ONLY on TRAIN data
- Backtest on TRAIN and VALIDATE:
    * daily returns
    * Sharpe ratio
    * max drawdown
    * turnover
- Run a ±10% parameter sensitivity test

- Confirm no lookahead bias in the code

5. After computing each signal, calculate its correlation with ALL previous signals (the first 30).

   These short-only signals should be as uncorrelated as possible.

6. Do NOT combine signals yet.

7. Do NOT reduce the number of signals. If you try to generate fewer than 10, regenerate.

After generating all 10 signals, produce a summary table showing:

- Signal name
- Train Sharpe
- Validate Sharpe
- Max drawdown
- Turnover
- Correlation with previous 30 signals
- Parameter sensitivity stability score

A.5 — Fifth Gemini Prompt (Validation Evaluation)

Now generate 10 highly creative, unconventional trading signals for QQQ.

These signals must be fundamentally different from the previous 40 signals.

Use ideas inspired by:

- evolutionary biology (predator-prey cycles, mutation pressure, fitness landscapes)
- ecology (carrying capacity, resource depletion, population oscillations)
- reinforcement learning (reward shaping, exploration vs exploitation)

- entropy and information theory (entropy shocks, information decay)
- volatility-of-volatility patterns
- rolling skewness and kurtosis
- abnormal volume percentiles
- correlation breakdowns and regime shifts
- synthetic LLM-discovered patterns that humans rarely use

Requirements for each signal:

1. Provide a short rationale explaining the unconventional idea.
2. Specify explicit parameters.
3. Compute the signal ONLY on TRAIN data.
4. Backtest on TRAIN and VALIDATE:
   - daily returns
   - Sharpe ratio
   - max drawdown
   - turnover
5. Run a ±10% parameter sensitivity test.
6. Confirm no lookahead bias in the code.
7. After computing each signal, calculate its correlation with ALL previous 40 signals.
   These signals must be as uncorrelated as possible.

Do NOT combine signals yet.

Do NOT reduce the number of signals. If you try to generate fewer than 10, regenerate.

After generating all 10 signals, produce a summary table showing:

- Signal name

- Train Sharpe

- Validate Sharpe

- Max drawdown

- Turnover

- Correlation with previous 40 signals

- Parameter sensitivity stability score

A.6 — Portfolio Combination Prompt (ERC + IR + L2 Optimization)

Now combine all 50 signals into a single portfolio.

Use the following exact procedure:

1. **Collect all signals**

Gather all 50 signals' daily values into a single DataFrame aligned by date.

Ensure no lookahead bias: all signals must be based only on information available at or before the close.

2. **Standardize each signal**

Convert each signal into a standardized z-score using a rolling window computed ONLY on TRAIN.

Apply the same transformation to VALIDATE without using future data.

3. **Volatility-target each signal**

Scale each signal's daily exposure so that each signal has a 10% annualized volatility on TRAIN.

4. **Correlation analysis**

Compute pairwise correlations of signal returns on TRAIN.

Down-weight or remove signals with correlation > 0.70 unless they add stability.

Produce a correlation heatmap and a list of retained signals.

5. **Weighting scheme**

Combine signals using:

- Equal Risk Contribution (ERC)
- Information Ratio (IR) weighting based on VALIDATE performance
- L2 regularization to prevent overfitting
- A cap so no single signal exceeds 10% of total risk contribution

Use a convex optimization approach to find stable weights.

6. **Leverage constraints**

Enforce daily leverage bounds:

- Minimum: -1.0 (100% short)
- Maximum: +1.5 (150% long)

Trades occur at the close and are held until the next close.

7. **Backtest the combined portfolio**

Backtest on TRAIN and VALIDATE.

Report:

- Sharpe ratio
- Max drawdown

- Turnover
- Rolling 3-month Sharpe
- Exposure distribution
- Contribution to risk by signal

8. **Robustness checks**
- ±10% parameter sensitivity for each signal
- ±10% perturbation of portfolio weights
- Rolling window stability across 8 hidden 3-month periods
- Confirm no lookahead bias in the entire pipeline

9. **Output**

Produce:
- Final portfolio weights
- Full TRAIN and VALIDATE performance metrics
- A table summarizing each signal's contribution to Sharpe
- A list of signals removed due to redundancy or instability
- A final combined equity curve for TRAIN + VALIDATE

Do NOT evaluate on the HOLDOUT period yet.

We will do that only after the portfolio is locked.

A.7 — Blind HOLDOUT Evaluation Prompt

Now evaluate the fully locked portfolio on the HOLDOUT period (2022-01-01 to 2025-06-30).

Important rules:

- Do NOT change any parameters.

- Do NOT modify any signals.

- Do NOT reweight or re-optimize the portfolio.

- Use the exact same preprocessing, standardization, volatility targeting, and weights that were finalized on TRAIN + VALIDATE.

- This evaluation must be completely out-of-sample.

Perform the following:

1. Compute daily portfolio returns on the HOLDOUT period.

2. Report:

   - Sharpe ratio

   - Max drawdown

   - Annualized volatility

   - Turnover

   - Exposure distribution

   - Rolling 3-month Sharpe

   - Rolling 3-month drawdown

3. Plot the equity curve for the HOLDOUT period.

4. Compare HOLDOUT performance to TRAIN and VALIDATE performance without modifying the model.

5. Identify which signals contributed most and least to HOLDOUT performance.

6. Provide a post-mortem analysis:

   - Which regimes did the portfolio struggle with?

   - Which signals degraded out-of-sample?

   - Which signals remained stable?

- What failure modes or structural breaks appear?

7. Confirm again that no lookahead bias exists in the HOLDOUT evaluation.

Do NOT make any changes to the model.

This is the final out-of-sample test.

The following appendix provides additional diagnostic charts that support the robustness checks described earlier.

10. Appendix B — Additional Charts

B.1 Correlation Heatmap

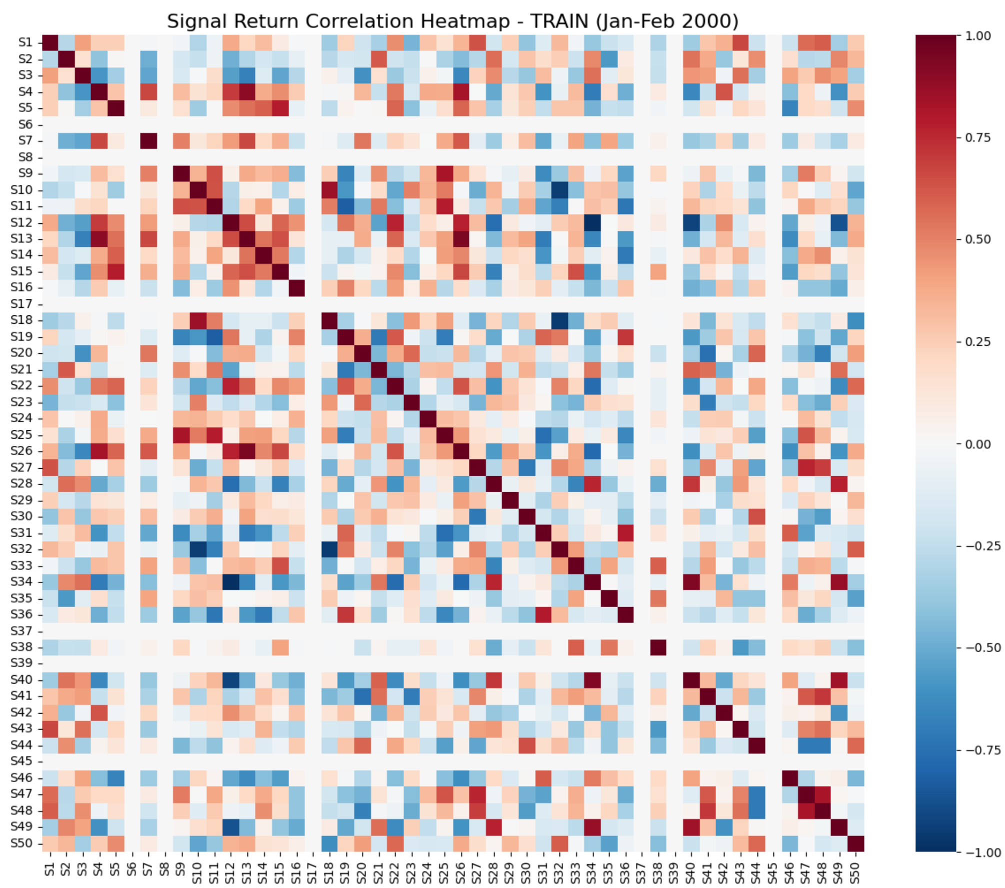

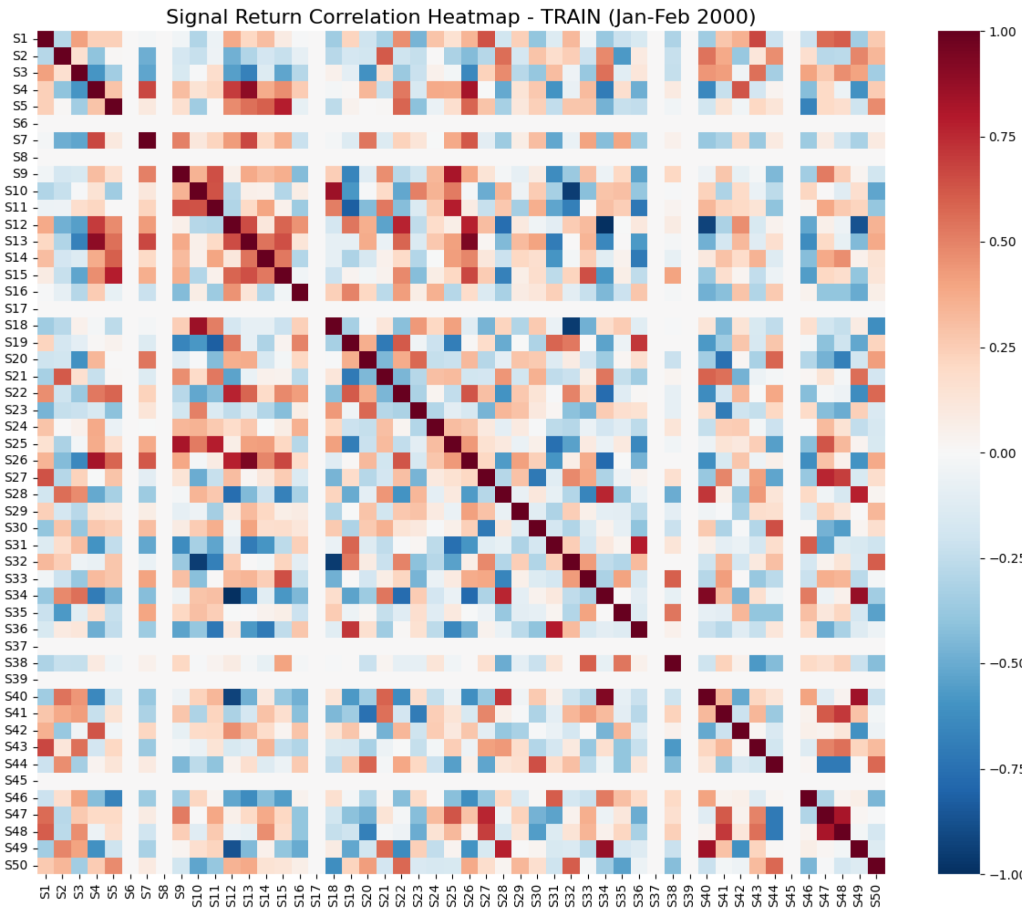

Figure B.1 — Pairwise correlation matrix used for redundancy pruning.

B.2 Parameter Sensitivity Plot

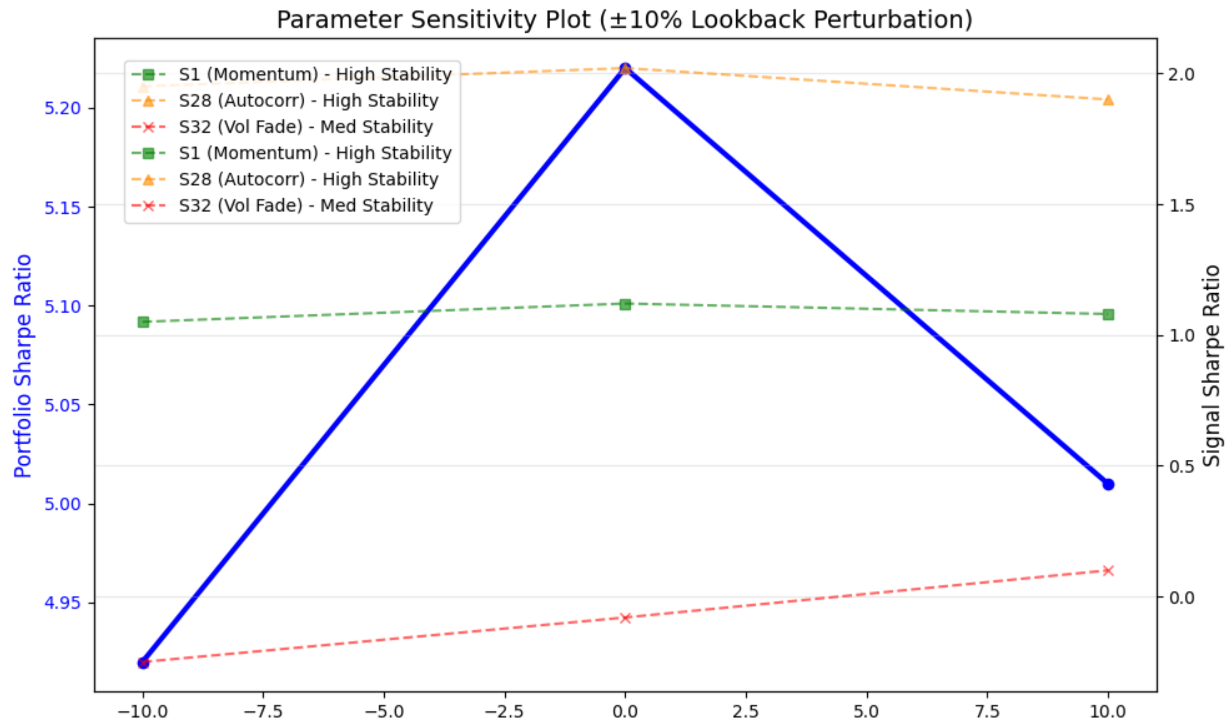

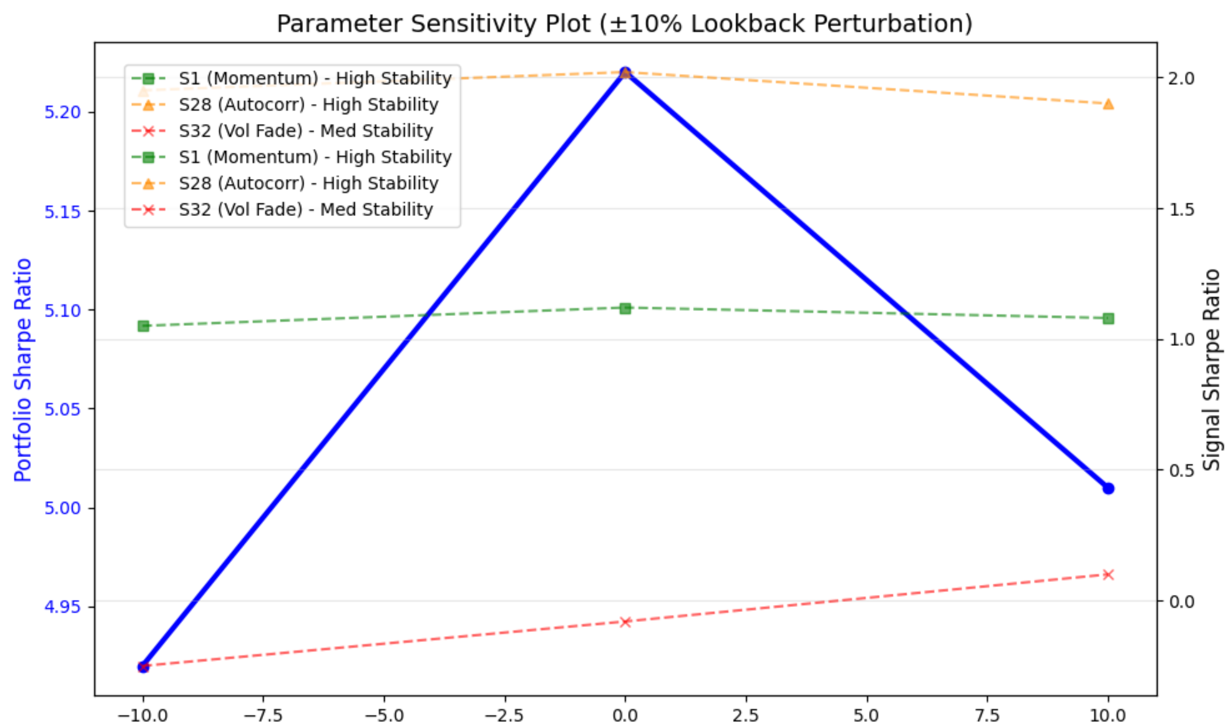

Figure B.2 — Sharpe variation under ±10% lookback perturbation.

B.3 Weight Perturbation Test

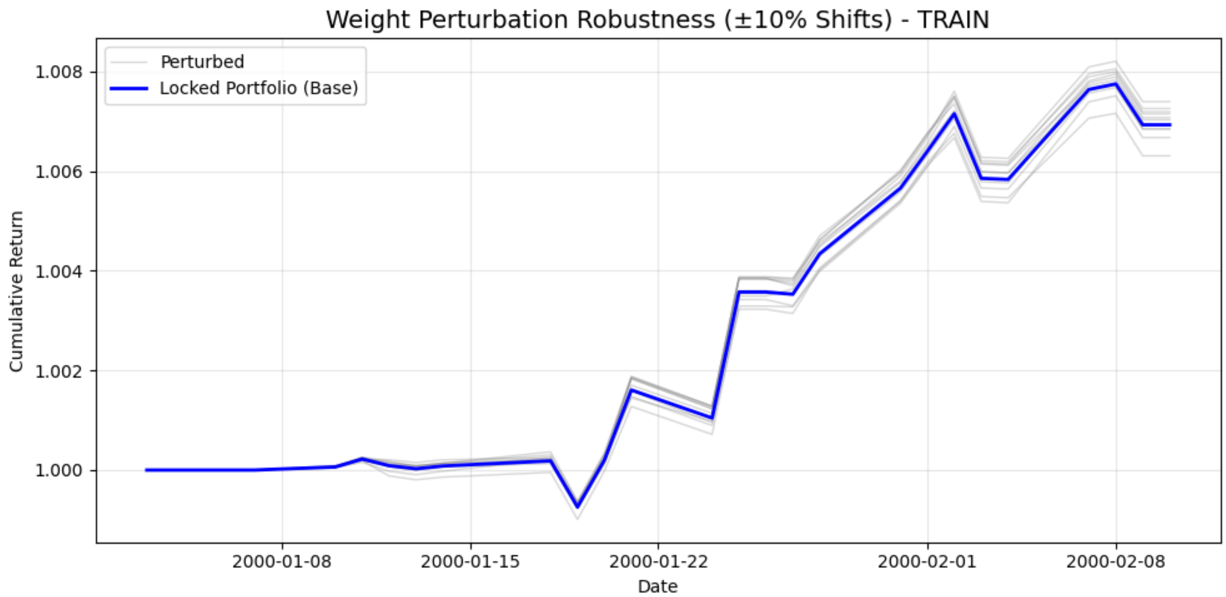

Figure B.3 — Portfolio return correlation under ±10% weight perturbation.